

High acquisition, low retention

A read on 91 months of orders — financial health, the retention problem underneath the growth, and where the next point of profit actually is.

● 91 months analyzed

SQL · Python · public dataset

Great top line, hiding a retention problem

\$2.5M

All-time completed revenue over
91 months

\$133K

Latest full month (Aug 2026) ·
1,535 orders

~52%


Average gross margin · healthy
unit economics

~96%

90-day churn rate · recently ~93%

Revenue compounds and it is profitable — but the business loses almost every buyer within 90 days. The growth is almost entirely **new** customers, not repeat behavior.

Revenue keeps compounding

**\$2.5M**

All-time revenue

~52%

Average gross margin

\$1.3M

All-time gross profit

~\$87

Latest-month AOV · 1,535 orders

Durable, roughly 30× growth since 2019 — and every dollar carries a healthy ~52% margin. The top line is not the problem.


Almost no one comes back

90-day churn stays extreme as revenue scales



Of customers active in a month, the share with no order in the next 90 days — averaging **~96%** (recently **~93%**) even as revenue grew **~30x**.

The month-0 → month-1 cliff




Of every 100 first-time buyers, only **~3** purchase again the next month — then it flatlines near 1%.

Growth is real, but it is rented, not owned: the business re-buys its revenue every month.

New customers carry the top line

New-customer revenue Returning-customer revenue



~9% → ~15%

Returning share of revenue: long-run average vs latest month

Returning revenue has climbed from near zero to ~15% — genuine progress, but four of every five dollars still come from someone buying for the first time.

Three moves, in priority order

01

Attack retention, not just the funnel

Stand up a lifecycle program — post-purchase onboarding and a day-30/60 win-back — aimed squarely at the month-0 → month-1 cliff. Track repeat-rate at 30/60/90 days as the north-star metric.

02

Grow returning-revenue share on purpose

It drifted to ~15% on its own — make it a target. Test loyalty and replenishment nudges on first-time buyers in the highest-margin categories, and measure *incremental* repeat purchases.

03

Instrument before pricing & shipping bets

Add order-level **shipping_fee** and promo flags, define the launch date, and reserve a holdout — turning a “free shipping over \$100” launch into a measurable difference-in-differences test.

The through-line: a few points of retention compound faster than the same dollars spent re-acquiring customers the business already had.

How this was measured

Scope & window

91 complete months, 2019-02 → 2026-08. The current partial month is excluded so growth and trend lines are not distorted by an unfinished month.

Churn is deliberately simple

The 90-day churn is calendar-anchored and binary, reported only for fully-observed months. In production I'd anchor the window to each customer's last order.

Definitions

Completed sale = status **Complete** and not returned. Revenue = Σ sale_price; margin uses inventory cost. New vs returning is keyed on each customer's first completed order.

Free shipping is a scenario, not data

thelook has no shipping-fee field, so that task builds the difference-in-differences *frame* you'd use to measure such a policy — rather than reporting a fabricated effect.

THANK YOU

Questions?

Happy to walk through the SQL, the notebook, or any assumption behind these numbers.

● Product Analyst · Data Challenge